

Overlapping construction plus a position-blind split: a curation defect in genomic sequence benchmarks and the model rankings it inverts

Rikhin Kavuru^{1,*}

¹Independent Researcher, United States

*Corresponding author. rikhinkavuru@gmail.com

Abstract

Motivation: Genomic benchmarks rarely control for near-duplicate sequences spanning their train/test split. A model can then score by recognizing memorized near-copies rather than generalizing, and no audit locates its origin.

Results: Near-duplicate leakage here requires two conditions together—an overlapping construction step and a split that ignores genomic position—rather than any property of the sequence type. A pre-registered prediction that short human regulatory tasks would be leaky failed: those GUE tasks are clean at full scale across seven independent partitions (eleven shipped tasks). The signature precedes any model: in Genomic Benchmarks the leaky positive class ships 77,421 intervals on 40,934 distinct coordinates, and a parameter-free coordinate screen separates seven of eight datasets. The defect is causal both ways: the omitted merge step drops the leak fraction from 0.390 to 0.012 and the random forest's gap from +0.165 to −0.008; imposing duplication manufactures both in a clean dataset. The headline drop separates into two effects that a single number hides: on `human_enhancers_ensembl` 9.0 of the memorizing forest's 16.4 points is evaluation-side inflation, measurable on the novel stratum with no retraining and no re-split, and the remaining 7.4 is a retraining penalty paid by that model alone. The benchmark then cannot rank models: the as-shipped split leaves its top three of nine learners within 0.0015, yet corrected they span 23 points—more than a difficulty-matched test set at the same accuracy produces (14), and on two clean datasets the same re-split expands spread not at all; a from-scratch CNN ranks fifth of ten as shipped and first corrected. Of the two leaky datasets only that one supports the ranking claim: `human_nontata_promoters` certifies leaky-but-correction-unsafe, its components being single-linkage chains, so its corrected delta is reported but not relied on.

Availability and implementation: Code, the report card, a near-duplicate-aware splitter and a certification tool are at <https://github.com/rihinkavuru/homology-leakage-genomic-benchmarks>.

Contact: rikhinkavuru@gmail.com

Supplementary information: Available at Bioinformatics Advances online.

Keywords data leakage, near-duplicate sequences, benchmark evaluation, genomics, model ranking

1. Introduction

Benchmark suites are the default proving grounds for new DNA sequence classification methods, and the Genomic Benchmarks collection (Grešová et al., 2023) is used in part because it is small and quick to run. The value of any such suite rests critically on one assumption among several: that its train/test split measures generalization rather than recall of the training data. Near-duplicate leakage—exact or near-identical sequences appearing on both sides of a split—breaks that assumption, because a model can then succeed by recognizing memorized near-copies, so reported accuracy overstates generalization.

The evidence for such leakage can sit in files distributed with the benchmark, before any model is trained. The positive class of `human_enhancers_ensembl` ships 77,421 genomic intervals on only 40,934 distinct coordinates: 36,487 coordinates appear exactly twice and 4,447 once, so 94.3% of positive rows sit on a duplicated coordinate, and because the shipped split is

random, 75.5% of test positives have their exact coordinate present in training. A near-default random forest (150 trees, `min_samples_leaf=1`) scores 0.860 on that split and 0.696 once the near-copies are removed; on the near-duplicate test sequences alone it scores 1.000, and on the rest 0.770—worse than the linear SVM it outranks. Four of the suite's seven binary datasets carry none of this, one is borderline under a length-robust metric, and two are leaky, so the defect is a property of particular datasets rather than of the suite.

Leakage is an established benchmarking hazard. Kapoor and Narayanan (2023) document it as a pervasive driver of the reproducibility crisis in ML-based science, inflating performance and, in extreme cases, changing the apparent ordering of methods. In genomics, hashFrag (Rafi et al., 2026) shows homologous sequences spanning cis-regulatory splits inflating neural-model performance, and reports that models rely on memorized associations wherever function is

conserved between homologs; standard protein-interaction splits leave most test cases with near-duplicate counterparts (Bushuiev et al., 2024); de-leaked test sets sharply reduce reported performance including for AlphaFold-Multimer (Kovtun et al., 2024); and on LIT-PCBA a trivial memorization baseline matches sophisticated methods because the benchmark is leaky (Huang et al., 2025). Evaluation design is itself a recognized genomics hazard (Whalen et al., 2022; Schreiber et al., 2020; Toneyan et al., 2022). Outside biology the ranking question has been posed most sharply: Barz and Denzler (2020) find near-duplicates of training images inside the CIFAR test sets, Northcutt et al. (2021) show test-set defects destabilizing the ranking of models evaluated on ten widely used benchmarks, and under controlled injection Zhou et al. (2023a) show benchmark leakage distorting model comparison, a hazard established at scale by work on train/test overlap, corpus deduplication and language-model contamination (Elangovan et al., 2021; Lee et al., 2022; Sainz et al., 2023; Magar and Schwartz, 2022; Deng et al., 2024). The usual remedy is a similarity-aware partition (Hobohm et al., 1992; Walsh et al., 2016; Teufel et al., 2023; Petti and Eddy, 2022; Li and Godzik, 2006; Steinegger and Söding, 2017).

What that literature lacks, and a small curated suite supplies, is a matched clean control inside the same collection and a construction step that can be edited. The open question is accordingly upstream of the remedy: not how to repair a leaky split, but where the leakage originates and whether it reorders which model wins. Two candidate answers are separable. Leakage may track what a dataset contains—sequence type, task, window length—in which case a practitioner can screen by task family. Or it may track how the dataset was assembled, in which case the screen belongs in coordinate space and the fix belongs in the curation pipeline. Distinguishing them requires a prediction committed in advance and an intervention on the construction step alone. Pretrained foundation models are outside this evidence: DNABERT-2 (Zhou et al., 2023b), HyenaDNA (Nguyen et al., 2023) and the Nucleotide Transformer (Dalla-Torre et al., 2024) all pretrain on the GRCh38 human reference—HyenaDNA on it alone, the other two as the human component of a multispecies corpus—and these test sequences lie in it (§4), so evaluating them would open a second, pretraining-based leakage channel no split can close.

Across the datasets audited here, leakage tracks how a dataset was assembled rather than what kind of sequence it contains, and the condition it requires is conjunctive: an overlapping construction step *and* a split that ignores genomic position (§4). Where similarity is instead a property of the sequences themselves—GUE’s *virus.covid*, whose nine SARS-CoV-2 variants differ by a handful of mutations—a leak fraction of 1.000 is biology and carries no information about curation (§4). We measure a coordinate-space signature on the shipped genomic intervals rather than the sequences, and a screen with no fitted parameter separates the suite 7/8 and calls four out-of-suite BEND partitions clean; we edit the construction step in both directions, applying the omitted merge step to a leaky dataset to abolish the leakage and its downstream consequences, and imposing it on a clean dataset to manufacture both; and we score a prediction committed before the target suite was examined, that GUE’s short human regulatory tasks would be

leaky, which failed—all eleven are clean at full scale, refuting the task-type hypothesis the prediction encoded. The downstream cost is a benchmark that cannot rank models, demonstrated on *human.enhancers.ensembl*, where the as-shipped split cannot separate its top three of nine learners—within 0.0015—yet those three span 23 accuracy points corrected, more than the 14 a difficulty-matched test set at the same operating point produces. That demonstration requires a memorization-prone model in the comparison set and a split whose construction leaves near-duplicates spanning it; both datasets on which it holds are short human regulatory sets, one fixed at 251bp and one variable at a 269bp median over a 2–573 range, so that regime remains a possible further necessary condition, and is not a sufficient one. We release a per-dataset report card, an aligner-free near-duplicate-aware splitter, and a certification tool returning *leaky-but-correction-unsafe* where whole-cluster correction is not safe.

2. Methods

2.1. Datasets

We study the seven binary Genomic Benchmarks datasets (Grešová et al., 2023) and additionally report the three-class *human.ensembl.regulatory* set; sizes, lengths and balance are in Supplementary Table S1. Their upstream provenance is pinned rather than assumed: per the curators’ own construction notebooks, *human.nontata.promoters* and *human.enhancers.cohn* derive from the Ensembl release-97 GRCh38 toplevel DNA FASTA, and *human.enhancers.ensembl* and *human.ocr.ensembl* from the release-100 Regulatory Build external-feature table over the release-100 GRCh38 primary assembly. Every upstream URL was re-resolved on 2026-07-20 (all HTTP 200; results/reference.provenance.csv), and archive hosts are recorded alongside the live ones—<https://jul2019.archive.ensembl.org> for release 97 and <https://apr2020.archive.ensembl.org> for release 100—because Ensembl release 116 (June 2026) is the last served on the current platform, with data moving to a new site from summer 2026, so the live paths in the notebooks are not durable citations. Balanced datasets are curated to 0.500 positive fraction in both train and test, so class balance is preprocessed, not natural. All censuses and leak fractions are full-scale; four accuracy deltas are not. The corrected-split deltas for *human.ocr.ensembl*, *demo.human.or.worm* and *human.enhancers.cohn*, and the clean-control roster arm of §3.6, run on 20,000-sequence subsamples of datasets whose report-card column is headed *n* (full). Their full-scale leak fractions are ≤ 0.012 (Jaccard) and ≤ 0.068 (containment) for the three binary sets, and subsampling can only *hide* leakage, so a subsample is admissible there to confirm cleanliness—never to establish it. The three-class set is bounded on the Jaccard side alone (0.005 at full scale): its full-scale containment was never computed, and since containment is precisely what moved *demo.coding.vs.intergenomic.seqs* from clean to borderline in Table 1, we state that gap rather than let the binary sets’ bound stand in for it. On a 20,000-sequence subsample its containment reads 0.014, which by the same one-way argument confirms nothing.

2.2. Coordinate-space construction signatures

Genomic Benchmarks distributes per-sequence intervals (`id,region,start,end,strand`) separately from the sequences its loader returns. We recover them offline and realign them to loader row order, which is deterministic (split, then sorted class directory, then sorted filename, the stem being the interval `id`); recovered interval width equals sequence length for 100% of rows in all eight datasets, a check informative only for the four variable-length ones. These coordinates are read on a separate code path from the k -mer clustering and are causally upstream of it, so they give non-circular corroboration rather than a restatement. Per class we compute self-overlap redundancy $R_{\text{self}} = 1 - n_{\text{merged}}/n_{\text{raw}}$ and the share of test intervals overlapping any *train* interval at $\geq t$ reciprocal overlap, $t \in \{\geq 1 \text{ bp}, 0.5, 0.9, 1.0\}$; overlap is strand-agnostic, strand being unrecorded for the two *demo* datasets. The screen thresholds the maximum over classes of the $\geq 50\%$ statistic at 0.1—the cut already used in sequence space—so it adds no fitted parameter.

The assembly is measured, not assumed, and the unsampled check is the stronger one. All 190,973 shipped intervals of the two leaky datasets name a primary chromosome and none an alternate contig, and—without sampling, so with no denominator to choose—915 of them (0.48%) end past a GRCh37 chromosome boundary while none ends past a GRCh38 one. A sampled sequence-level check agrees. For 49 test intervals stratified over both datasets, classes and strands we fetched the reference at the shipped coordinates from GRCh38 and GRCh37 via the UCSC REST API on 2026-07-20—hg38, the Dec. 2013 assembly, accession GCA.000001405.15, and hg19, the Feb. 2009 assembly, accession GCA.000001405.1—reading intervals as 0-based half-open and reverse-complementing minus-strand records: GRCh38 reproduces the shipped sequence in 49 of 49 cases and a 1-based reading in none of 49. The GRCh37 denominator is 48, not 49, for the same reason the boundary count gives: one sampled chrX interval runs past the GRCh37 chromosome end, so the API returns an error rather than a sequence and that row is unevaluable on GRCh37 by construction. On the 48 evaluable rows GRCh37 reproduces the shipped sequence once.

2.3. Construct-and-break manipulations

We intervene on the construction step alone. *Fix-a-leaky*: *human_enhancers_ensembl* receives the merge step its curators omitted—every duplicated coordinate collapsed to one representative row—with sequences, labels, model code and split ratio unchanged and negatives subsampled to restore the control's balance. *Break-a-clean*: *human_ocr_ensembl*, drawn from the merged Ensembl Regulatory Build and clean by every measure, is rebuilt as an un-duplicated assembly, duplicating positives until the same 94.3% share of positive rows sits on a multiplicity-2 coordinate as on *human_enhancers_ensembl* (requiring 89.1% of them, since duplicating a fraction r leaves $2r/(1+r)$ of rows on a duplicated coordinate), then downsampling to the control's exact size and class balance. Each manipulation is paired with an unmanipulated control at the same split seed. In both manipulated datasets the duplicated

coordinates sit in the positive class, so both interventions perturb balance and size unless corrected; balance-matched arms are primary, and the break-a-clean pair is matched on size as well while the fix-a-leaky pair cannot be, since size there cannot be equalised without discarding the rows the intervention is defined to remove. A ten-dose version sweeps the imposed duplication rate (Supplementary §S2).

2.4. Near-duplicate measurement and re-split

We measure train/test overlap by the exact Jaccard similarity of k -mer sets ($k = 8$) over all test–train pairs via sparse binary matrix products; the leakage fraction is the share of test sequences whose maximum similarity to any training sequence exceeds a threshold (0.5, 0.7, 0.9). A k -mer Jaccard is bounded by the length ratio, so it undercounts leakage on datasets with disparate lengths—across *human_enhancers_ensembl*'s test–train pairs, 61% have their Jaccard held below 0.7 by length ratio alone—so we also report the length-robust *containment* index $|A \cap B| / \min(|A|, |B|)$. Jaccard leak fractions are lower bounds; unfloored containment runs the other way and is an *upper* bound on variable-length data, since a very short sequence is contained in almost anything. Calibrated on synthetic sequences with known ground truth, the 8-mer Jaccard at 0.7 detects an exact copy, one point mutation and a window shifted by up to 10% of its length at every length we use, but at 70 bp it misses two point mutations 70% of the time where at ≥ 300 bp it tolerates five: a near-exact-duplicate detector on short sequences, a homology detector on long ones, so threshold choice should follow length (Supplementary §S4). We separately census exact byte-identical duplicates on a code path using no k -mers: *human_enhancers_ensembl* has 11,774 such test sequences (0.380 of its test set), whereas *human_nontata_promoters* has 0 exact duplicates yet 0.225 near-duplicates at Jaccard ≥ 0.9 —the more homology-like but statistically weaker case.

To correct leakage we connect sequences whose Jaccard exceeds 0.7, take connected components as clusters, and assign whole clusters to train or test, preserving the original ratio and balance and guaranteeing zero residual cross-split similarity above the threshold. Sweeping the threshold, the drop is flat on *human_enhancers_ensembl* (0.156/0.156/0.157 at 0.5/0.7/0.9; 0.158 under largest-component removal) but scales strongly on *human_nontata_promoters* (0.205/0.121/0.006), since 0.9 leaves its large [0.5, 0.9] band uncorrected. This is an aligner-free instance of GraphPart's whole-cluster partitioning (Teufel et al., 2023); an MMseqs2 arm is in Supplementary §S3.

2.5. Models, decision rule and pre-registration

Sequences are represented by overlapping k -mer counts, $k \in \{4, 6\}$. Nine classifiers span *memorization propensity* (equivalently, effective regularization) rather than raw capacity: logistic regression, a linear SVM (LinearSVC), a random forest (150 trees at scikit-learn defaults, `min_samples_leaf=1`), HistGradientBoosting, extremely randomized trees, 1- and 15-nearest neighbours, a multilayer perceptron (MLP) and

Gaussian naive Bayes; the first four constitute the four-model comparison, all nine the roster. Scale-sensitive learners receive L2 normalization, tree ensembles raw counts. All are trained from scratch on each split, so none carries a pretraining confound. Accuracies use `.predict()` (`argmax`); the headline drop is the original accuracy minus the mean corrected accuracy over five re-split seeds. Parameters and roster rationale: Supplementary §S1. An inversion criterion built for four models is too permissive for nine—"the top model changed by more than a point" fires on a *clean* control from repartitioning alone—so we additionally require that the two models exchanging the top slot each lead on their own split by a *paired* cluster-bootstrap difference excluding zero, and we report the swap margin separately from the as-shipped first-to-second margin, which can be a tie when the former is large.

The deep-model probe is a from-scratch 1D residual convolutional network in the Basset (Kelley et al., 2016) / DeepSEA (Zhou and Troyanskaya, 2015) lineage (one-hot input, no *k*-mer features), trained on each dataset's training split only. We pre-specified, in a document committed alongside the results it predicts, the narrowed claim, a dropout×weight-decay grid, and a refutation condition: if the standard-practice regularized cell (dropout 0.3, weight decay 10^{-4}) has a cluster-bootstrap drop CI including 0 on *both* datasets, the strong "any expressive learner is inflated" reading is refuted and reported as a limitation. This is a binding stated condition, not an externally timestamped registration; it was committed in the same commit as the experiment code and so carries no priority evidence at all, unlike the cross-suite registration below. Commit hashes and timestamps for both, with the `git log` output that establishes the ordering and an explicit statement of what author-controlled timestamps can and cannot show, are published as `PREREGISTRATION.md` in the code repository.

Six predictions were committed in code, inside the module that would test each, before that module was run, and all are scored in the Results whether or not they held: P1 1-nearest-neighbour shows the largest graded memorization gap, registered against the *raw* form of that statistic—the prevalence-corrected form is exploratory here and is registered forward for the next dataset, not counted backward (§3.7); P2 it shows the largest split-induced drop; P3 the four memorization-prone learners (1- and 15-nearest neighbours, the forest, extremely randomized trees) outrank the two linear models as shipped and fall below them after correction; P4 on a clean control no model's drop interval excludes zero; P5 cluster-grouped cross-validation selects a more regularized forest than naive cross-validation; P6 a ranking inversion occurs at every imposed leak fraction above the diagnostic threshold of Supplementary §S2 and at none below. Separately, before that suite was examined, we registered for GUE (Zhou et al., 2023b) that its short fixed-length human regulatory tasks would be leaky and its multi-species tasks clean (§3.12).

The *graded memorization gap* is $\text{acc}[\geq 0.9] - \text{acc}[\leq 0.5]$. Raw, it is confounded: the strata differ sharply in class composition, in opposite directions on the two leaky datasets, and a constant classifier that memorizes nothing scores +0.803 on *human_enhancers_ensembl* and −0.980 on *human_nontata_promoters*. We therefore report the gap from *balanced* accuracy within each stratum, for which a constant

classifier scores exactly zero, alongside the raw gap so the confound stays visible (Supplementary §S1).

2.6. Controls, cross-suite census and intervals

Chromosome holdout and regularization.

Using the recovered coordinates we assign whole chromosomes to the test side in a single greedy pass until the target test fraction is reached (nine chromosomes for *human_enhancers_ensembl*, twelve for *human_nontata_promoters*; there is no cross-validation loop). This field-standard control is deployment-relevant but *not* leakage-free: residual leakage at Jaccard 0.7 is 0.8% and 0.02%, against the exact 0 the near-duplicate-aware split guarantees. Separately we sweep, for the random forest only, `min_samples_leaf` ∈ {1, 5, 20, 50, 100, 200, 500} and `max_depth` ∈ {None, 16, 8, 4} at full scale—the largest settings exceed the largest cluster, so no leaf can be a pure duplicate block—and cross-validate the same leaf grid on the as-shipped *training* set under ordinary stratified and cluster-grouped 5-fold.

Cross-suite census.

The census is applied unchanged to the Nucleotide Transformer downstream tasks (Dalla-Torre et al., 2024) and to GUE (Zhou et al., 2023b); the full pipeline is then run on the tasks the first flags, with a clean task as control. Collapsing two nestings leaves eleven independent tasks of the thirteen shipped, and we report independent-task counts throughout; the three splice tasks are not nested and count separately, and we draw no before-and-after inference between the original and revised releases, which are not the same data recaptured. One deviation affects the anchoring task: *enhancers* ships only 400 test sequences, too few for a usable interval, so for that task alone we pool train and test and draw a fresh stratified 80/20 split before running both arms. Pooling redistributes the duplicate pairs, so the re-split leak fraction ($\varphi = 0.099$) is not the same quantity as the 25% byte-identical figure the as-shipped split carries (Supplementary §S1).

Injected multiclass construction.

The only multiclass set is clean, so we answer the question by construction, injecting label-carrying near-duplicate train→test copies. Here *f* is the fraction of *training* rows copied and appended to the test set, so realised test-set leak fractions are $\varphi = 0.29, 0.44$ and 0.62 at $f = 0.1, 0.2, 0.4$, not *f*. The arm runs on a 20,000-sequence subsample; the dose is imposed rather than measured, so scale is not load-bearing.

Intervals and reproducibility.

Because near-duplicates violate the independence a sample-wise bootstrap assumes, we resample whole within-test near-duplicate components as blocks on *both* arms, quantify the design effect using Kish's effective block size rather than the mean (these distributions are heavy-tailed and the mean-size convention understates the factor roughly twofold), report it beside the cluster-versus-sample variance ratio the bootstrap measures directly, cross-check against a Liang–Zeger cluster-robust standard error, and fold the five re-split-seed variance into a combined-source interval; Benjamini–Hochberg at $q = 0.05$

is applied over a declared family of 15 deltas (Benjamini and Hochberg, 1995), whose membership rule is stated rather than left to inference: every per-dataset corrected-split accuracy delta in Table 1 and §3.5, for the two pre-registered model families, on all eight datasets. The cross-suite ranking grid is corrected as its own family of 16, and because two families corrected separately is the more permissive choice we report the pooled sensitivity analysis in §3.13. Two properties of these p -values are worth stating: they are recovered from percentile intervals under a normal reference, $SE \approx (hi - lo)/2z$, because the bootstrap draws were not retained, and they are therefore adequate for a verdict question only where deltas sit far from the cut—which is why any delta landing within a factor of two of q is flagged as unresolvable rather than called. Reported 95% CIs are percentile bootstrap intervals over 10,000 draws with the resampling unit stated per number; the three combined-source intervals are instead Wald, $\delta \pm z\sqrt{s_{boot}^2 + s_{seed}^2}$, and labelled where they appear. Models are not refit while the test set is resampled, so every interval is conditional on the fitted models. The classical pipeline is deterministic, seeded and CPU-only; the CNN is not, since 12 of 20 committed grid cells carry dropout > 0 and dropout masks are drawn from the device RNG, so CNN results are device-dependent within the five-seed tolerance we report. Because *genomic_benchmarks* 1.0.0 re-extracts on every call we cache locally and pin the cached bytes by SHA-256, recording row counts, extraction timestamps and both file and content digests for all nine cached combinations (Supplementary §S1.8); audits against a different release should re-derive the report card and will see the content digests move.

3. Results

3.1. Leakage is dataset-specific, and one clean verdict moved

Only *human_enhancers_ensembl* and *human_nontata_promoters* exceed a 10% near-duplicate test fraction at Jaccard 0.7 (0.384, 0.406); the three-class set is clean at 0.005 (Table 1). Re-measuring at full scale with containment reveals substantially more leakage on *human_enhancers_ensembl* (0.384 $\rightarrow$ 0.845) and flags one dataset the length-blind metric misses: *demo_coding-vs_intergenomic_seqs*, clean under Jaccard (0.078), crosses the cut under containment (0.131), so we relabel it *borderline*. The remaining four stay well below under both metrics (largest containment 0.068). Two datasets are leaky—one of them, *human_nontata_promoters*, certified *leaky-but-correction-unsafe* because its largest-cluster cohesion is 0.084, so whole-cluster correction would strip genuine signal along with the leakage—one *borderline*, four clean. A clean verdict moved, so both metrics should be reported.

3.2. The cause is a curation defect visible in the coordinates

Of the 77,421 positive intervals in *human_enhancers_ensembl*, only 40,934 are distinct: 36,487 coordinates ship exactly twice and 4,447 once, so 94.3% of positive rows sit on a multiplicity-2 coordinate (Fig. 1A). That coordinate redundancy is the sequence redundancy to a close approximation, and the two

counts coincide to 0.7%. In coordinate space, 11,694 test positives ship a coordinate that also occurs among the training positives—a direct count on the shipped split, and 0.15% above the 11,676 a class-stratified uniform split predicts, $36,487 \times 2 \times \frac{15,485}{77,421} \times \frac{61,936}{77,420}$, one duplicated coordinate contributing such a row exactly when one of its two rows lands in test. In sequence space, on a code path using no coordinates, 11,774 test sequences are byte-identical to a training sequence. The agreement is between two independent counts; per-pair identity was not checked on all 36,487 pairs.

Thresholding the maximum-over-classes cross-split coordinate overlap at the same 0.1 cut used in sequence space separates the suite 7/8, the two leaky datasets reading 0.755 and 0.992 against 0.073 and below for the rest (Fig. 1B). The single miss is *demo_coding-vs_intergenomic_seqs*, containment-borderline but coordinate-clean: its leakage is coding-paralog homology and its *coding_seqs* class is keyed by transcript accession, one region per row, so its coordinate statistics are structurally zero rather than measured; coordinates may escalate a verdict and never clear one. Two refinements were forced by the data: redundancy is not always in the positive class, *human_nontata_promoters* carrying essentially all of its own in the *negative* class ($R_{self} = 0.961$ against 0.047); and mere adjacency over-calls, *drosophila_enhancers_stark* tiling at a 2142 bp median with base-pair-scale touches, 0.406 at ≥ 1 bp but 0.036 at $\geq 50\%$ reciprocal overlap.

A comparison internal to Ensembl ties the defect to a missing deduplication step: *human_ocr_ensembl* and *human_ensembl_regulatory*, from the release-100 Regulatory Build's merged "MultiCell" segmentation (Zerbino et al., 2015), sit at 0.000–0.010 self-overlap where the un-deduplicated *human_enhancers_ensembl* sits at 0.471. That is an association, not a cause—the upstream resources differ, so it is *not* one pipeline run with and without a merge step. Classifying all eight datasets by documented positive-class construction reproduces every verdict (Table 2). That classification is now anchored in the curators' own construction notebooks rather than inferred from our statistics: the two Ensembl datasets are separated at the configuration level, *human_enhancers_ensembl* requesting Ensembl's *external_feature* table and *human_ocr_ensembl* its *regulatory_feature* table, and the omitted merge is a documented absence—searching all four notebooks and the external scraper they delegate to, no step merges or removes overlapping intervals at any layer. Two qualifications keep that honest. The Ensembl notebooks contain no construction logic at all, deferring to a scraper we had to fetch separately, so the absence is established over both layers rather than over the notebook alone; and *fasta2loc* keys its results by sequence, so *exact* duplicates collapse incidentally, which makes the correct claim "no merge or overlap-removal step" rather than "no deduplication of any kind" (*results/gb_notebook_audit.csv*). The construction categories are therefore documented ex ante; only the mapping from category to verdict is ours, and the causal claim still rests on §3.3. The causal claim rests on §3.3.

Table 1. Near-duplicate leakage report card. **Verdict** is LEAKY if the full-scale near-duplicate test fraction exceeds 0.1 at Jaccard 0.7; a “clean” verdict means only *no forward-strand near-duplicate leakage detected*, subject to the disclosed length-pair cap—it is not a certificate of homology-freedom. “leak@0.7 (Jac / con)” gives the leak fraction under the 8-mer Jaccard and the length-robust containment index; “Acc. lost” is the best-model accuracy change under the near-duplicate-aware re-split (positive = accuracy lost). Leak fractions and censuses are full-scale on every row, but four of these deltas—*human_ocr_ensembl*, *demo_human_or_worm*, *human_enhancers_cohn* and the three-class *human_ensembl_regulatory*—are computed on 20,000-sequence subsamples of the datasets whose *n* column is headed *full*. Subsampling can only hide leakage, never create it, so a subsampled delta is admissible here to *confirm* cleanliness and never to establish it (§2.1). *demo_coding_vs_intergenomic_seqs* is *not* in that group: its delta is full-scale at *n* = 100,000. [†]*demo_coding_vs_intergenomic_seqs* is clean under Jaccard but borderline under containment. [‡]*leaky-but-correction-unsafe*: the verdict the certification tool returns for a LEAKY dataset whose largest-cluster cohesion falls below 0.5 (*human_nontata_promoters*, 0.084). Its near-duplicate components are single-linkage chains rather than sets of mutual near-copies, so whole-cluster re-splitting quarantines related but non-duplicate sequences and removes genuine signal along with the leakage. The leakage verdict stands; it is the *corrected* delta on that row that cannot be read as a clean measurement, and the novel-stratum accuracy on the as-shipped split should be reported in its place.

Dataset	<i>n</i> (full)	leak@0.7 (Jac)	leak@0.7 (con)	Verdict	Acc. lost	RF rank	Reorders under chrom.-holdout?
human_nontata_promoters	36,131	0.406	0.445	LEAKY [‡]	+0.118	1 → 3	no (RF stays rank 1)
human_enhancers_ensembl	154,842	0.384	0.845	LEAKY	+0.165	1 → 4	yes (RF → rank 4)
demo_coding_vs_intergenomic_seqs	100,000	0.078	0.131	borderline [†]	−0.001	4 → 4	n/a
human_ocr_ensembl	174,756	0.010	0.068	clean	+0.004	4 → 4	n/a
demo_human_or_worm	100,000	0.012	0.042	clean	−0.004	4 → 4	n/a
drosophila_enhancers_stark	6,914	0.016	0.033	clean	+0.012	2 → 3	n/a
human_enhancers_cohn	27,791	0.001	0.005	clean	−0.004	4 → 4	n/a
human_ensembl_regulatory (3-class)	289,061	0.005	—	clean	−0.010	n/a	n/a




Figure 1 The construction defect in the shipped coordinates, measured before any model is trained. (A) Distinct positive-class coordinates by the number of rows sharing them, for the leaky *human_enhancers_ensembl* and the clean *human_ocr_ensembl*; 36,487 of the former’s coordinates carry two rows each, 94.3% of its positive rows. (B) Fraction of test intervals overlapping any training interval at $\geq 50\%$ reciprocal overlap, taken as the maximum over classes, for the eight Genomic Benchmarks datasets and the four BEND partitions (green); the dashed rule at 0.1 is the LEAKY cut used throughout, carried over unchanged from sequence space, so the screen adds no fitted parameter. A measured 0.000 is drawn as an *open* hairline stub and labelled as such, so that it is distinguishable both from a small non-zero bar and from a missing one; four partitions read exactly zero. Overlap is strand-agnostic. *demo_coding_vs_intergenomic_seqs* is the one miss: it is borderline in sequence space under containment, but one of its classes is keyed by transcript accession rather than by interval, so its 0.006 is structurally zero rather than measured.

3.3. Editing the construction step alone creates and cures the reordering

Editing the construction step alone abolishes the syndrome in a leaky dataset and manufactures it in a clean one, turning the correlational coordinate evidence into a causal claim in both directions (Table 3). Applying the omitted merge step

to *human_enhancers_ensembl*, at matched balance, collapses its leak fraction from 0.390 to 0.012 and the random-forest inflation from +0.165 [0.157, 0.173] to −0.008 [−0.019, 0.003]—a null—and the reordering disappears, the forest sitting at rank 4 before and after correction. Imposing the same defect on the merged, clean *human_ocr_ensembl*, at identical size and balance, raises its leak fraction from 0.004 to 0.204 and its

Table 2. Positive-class construction step against measured verdict. **Retrospective:** categories were assigned with the verdicts already known, so the table has zero free parameters but is not a prospective test. The prospective coordinate call is BEND (§3.4); the prospective sequence-space call is GUE (§3.12), which the prediction got wrong. [‡]borderline under containment only.

Dataset	Construction of positives	Merge step	Verdict
h. enhancers_ensembl	un-dedup. assembly	none	LEAKY
h. nontata_promoters	fixed window, recurring loci	none	LEAKY
h. ocr_ensembl	merged segmentation	MultiCell	clean
h. ensembl_regulatory	merged segmentation	MultiCell	clean
h. enhancers_cohn	curated, deduplicated	curation	clean
d. enhancers_stark	curated, deduplicated	curation	clean
demo_coding_vs_interg.	random genome draw	disjoint draw	clean [‡]
demo_human_or_worm	random genome draw	disjoint draw	clean

Table 3. Construct-and-break manipulations. Each intervention is paired with an unmanipulated control run of the identical pipeline at the same split seed and matched class balance; the break-a-clean pair is matched on size as well. Control rows are fresh seed-0 random splits, not the shipped partition. “Leak” is the near-duplicate test fraction at Jaccard 0.7; “Rank” is the random forest before→after correction, of four models. Intervals are two-arm cluster bootstraps over whole within-test components; * marks intervals excluding zero, and control and intervention intervals are disjoint in both directions.

	Condition	<i>n</i>	Leak	RF drop [95% CI]	Rank
Fix	ensembl control	154,842	0.390	+0.165 [0.157, 0.173]*	1 → 4
	dup. collapsed	81,868	0.012	−0.008 [−0.019, 0.003]	4 → 4
Break	ocr control	20,000	0.004	+0.002 [−0.019, 0.025]	4 → 4
	pos. duplicated	20,000	0.204	+0.105 [0.083, 0.126]*	1 → 4

inflation from +0.002 [−0.019, 0.025] to +0.105 [0.083, 0.126], and manufactures the reordering: the forest is promoted to rank 1 on the contaminated split and demoted to rank 4 again by correction—the syndrome of §3.6, produced on demand. The effect is larger under matching than without it (+0.063 → +0.105), so the confounds were suppressing rather than creating it.

Duplicated rows scattered across a random split are leaky by construction; what the manipulation adds is that the full downstream syndrome, including the memorizer’s promotion to rank 1, follows from the construction step alone. Size matching removes some of the duplicate partners it has just introduced, so the imposed leak fraction (0.204) is about half the undownsampling value (0.498) and the estimate is conservative. A ten-dose sweep on the same clean dataset places the first ranking inversion at an imposed leak fraction of 0.106 and finds none below it, across eight doses beyond the two arms of Table 3 (Fig. 2; Supplementary §S2). Editing one curation step therefore both removes and installs the reordering, at a single fit per condition read only between arms.

3.4. An out-of-suite coordinate census

Because the screen needs only shipped coordinates it runs outside the suite, at lengths far beyond anything audited here




Figure 2 Ten doses of the same construction defect, imposed on the clean *human_ocr_ensembl*. Random-forest accuracy drop under near-duplicate-aware correction against the realised test-set leak fraction φ at Jaccard 0.9; the shaded band is the two-arm cluster-bootstrap 95% CI over whole within-test near-duplicate components, and each dose is a single fit at one split seed. Filled points are doses at which the four-model ranking inverts, open points doses at which it does not. Circled points are the two arms of Table 3 re-reported rather than fresh trials, so the sweep supplies eight additional doses. The dashed rule marks the first inversion, the dotted rule the diagnostic threshold φ^* computed at that dose.

in sequence space. Applied unchanged to the four supervised tasks of BEND (Marin et al., 2024), whose BED releases carry their own split assignments, it returns 0.000 at $\geq 50\%$ reciprocal overlap on chromatin accessibility (1,410,554 train against 372,153 test intervals), 0.000 on histone modification (433,861 against 120,567) and 0.000 on enhancer annotation in all ten folds—the last at 100,096 bp per sample (Fig. 1B). Gene finding, the one task partitioned at 80% identity rather than by chromosome, reproduces the *drosophila_enhancers_stark* pattern: 0.183 of its 597 test intervals touch a training interval at ≥ 1 bp, yet only 0.002 reach $\geq 50\%$ and none 90%. This is the only prospective construction-based call in the paper, and it is near-tautological on the three chromosome-split partitions. Together with the 7/8 in-suite separation, the screen therefore reads construction defects across four orders of magnitude of interval length without being refitted to the suite it was developed on.

3.5. A near-duplicate-aware re-split lowers accuracy only on leaky datasets

Under the re-split the best model’s accuracy drops sharply on the leaky datasets and nowhere else: 16.5 points on *human_enhancers_ensembl* ($0.8596 \rightarrow 0.6951 \pm 0.0019$ over five re-split seeds; combined-source Wald CI [15.4, 17.4]) and 11.8 on *human_nontata_promoters* ($0.9284 \rightarrow 0.8101 \pm 0.0055$). For that second dataset point estimate and interval are computed on different estimands. The five-seed mean

gives the 11.8-point figure; the combined-source interval—which additionally folds in the five re-split-seed variance, exactly as the *human_enhancers_ensembl* figure above does—is computed on the seed-0 partition, where the delta is 10.8 points, giving [8.8, 12.9]; the cluster bootstrap alone gives [9.2, 12.6]. Every clean-dataset delta includes zero for the memorizing forest as well as the linear models, the largest being *drosophila_enhancers_stark* at +0.024, cluster-bootstrap CI [−0.007, 0.054]. The negative control confirms causation: a random re-split at matched size and balance leaves the leaky-dataset accuracies unchanged (best-model deltas −0.001 and −0.002, point estimates, two orders of magnitude below 0.118 and 0.165; Supplementary Fig. S1), so only removing near-duplicates, not re-splitting per se, lowers the score. Corrected accuracy is stable across five re-split seeds and, independently, five forest training seeds (0.698 ± 0.002 ; 0.821 ± 0.003), the forest demoted on every seed. That de-leaking lowers apparent accuracy is established across domains (Kapoor and Narayanan, 2023; Rafi et al., 2026; Bushuiev et al., 2024); what the within-suite version adds is the matched clean control, on which the same procedure moves nothing.

3.6. The consequence: the benchmark cannot rank models

The four-model comparison, and where the fall comes from.

On *human_enhancers_ensembl* the best model *changes*. In the four-model comparison the near-default random forest, ranked first as shipped (0.860), falls to last after correction (0.696) and *LinearSVC* becomes best (0.798; Fig. 3). This holds under accuracy, AUROC (0.805 against 0.870) and F_1 (0.635 against 0.796), so it is not a threshold artifact; the bootstrap probability that the forest ranks first goes from 1.0 to 0.0; and it reproduces under the chromosome-holdout control (§3.13). The fall decomposes into two effects that a single number hides, and we separate them with a third arm: one fit on the as-shipped training set, scored twice—on the whole as-shipped test set and on its *novel* stratum alone (Table 4). **Evaluation-side inflation**, what the benchmark overstates about a model that already exists, needs no re-splitting and no retraining and accounts for 8.98 of the forest's 16.4 points ($0.8596 \rightarrow 0.7698$, cluster-bootstrap [8.7, 9.3]). **Split-side effect**, the retraining penalty from stripping near-duplicate partners out of training too, accounts for the remaining 7.41 ($0.7698 \rightarrow 0.6957$). Both are large, and the split-side half is not a measurement error at all. On *human_nontata_promoters* the same decomposition gives 4.91 and 6.64 points.

Two features of this decomposition matter more than its headline. First, the split side is where the memorizer is singled out: on *human_enhancers_ensembl* it is the *only* model of the four that pays a retraining penalty, the three others gaining slightly when refit on de-leaked data (−0.013, −0.013, −0.009), so the ranking change is not a rising tide. Second, the novel stratum is not class-balanced like the full test set (0.197 positive against 0.500; 0.982 against 0.544 on *human_nontata_promoters*), so the raw evaluation-side figure is prevalence-confounded in the same way §3.7's raw gap is—a constant classifier scores +0.303 and −0.438 on it for free. From balanced accuracy, for which a constant classifier scores zero,

the forest's evaluation-side inflation is +0.1789 and +0.1452, and the linear models' +0.0320 to +0.0499. Separately from either effect, the ranking on the novel stratum has *already* changed under the as-shipped fit, the forest scoring 0.770 against *LinearSVC*'s 0.782—a margin of 1.2 points, which is a different quantity from the 8.98 above and was previously reported in a way that invited the two to be confused.

Nine learners: a tie as shipped, 23 points corrected.

Widening from four learners to nine sharpens the statement (Supplementary Table S2). The leaky benchmark cannot separate its top three at all: the MLP (0.8736), 1-nearest-neighbour (0.8729) and extremely randomized trees (0.8721) finish within 0.0015, and the paired interval on the top-two margin, [−0.005, 0.006], includes zero. On the corrected split those three span 23 accuracy points (0.768, 0.537, 0.628). The benchmark is not merely crowning the wrong model; it is separating models it ranks as equals. Two controls bound how much of that separation is real, and they do not agree, so we report both and weaken the claim accordingly.

The first is a clean-dataset control, since the registered P4 governs *drops* rather than spread. Running the same nine-model roster and the same whole-cluster re-split on two clean datasets, top-three spread does not expand: on *human_enhancers_cohn* it goes 0.0218 [0.0120, 0.0328] as shipped to 0.0210 [0.0102, 0.0318] corrected, and on *demo_human_or_worm* it *contracts*, 0.0086 [0.0050, 0.0134] to 0.0056 [0.0026, 0.0104]. Re-splitting alone therefore does not manufacture spread.

The second control is harsher and is the one that constrains the claim. Top- k spread also expands mechanically as an operating point falls away from a ceiling, independently of whether the underlying differences are real. Re-weighting the *as-shipped* test set on a held-out difficulty panel until the same three models average the corrected split's accuracy (0.6951), and changing nothing else—no re-split, no retraining—their spread expands from 0.0015 to 0.1424 [0.1373, 0.1483]. That is most of the way to the corrected split's 0.2314. Difficulty alone therefore accounts for the majority of the expansion, and only the excess above it—roughly 0.09 of the 0.23—is attributable to de-leaking. We accordingly claim that correcting the split separates models the as-shipped benchmark reports as tied, by more than a difficulty-matched test set of the same accuracy would, and not that the full 23 points measures real model quality. The matched draw is also slightly *more* near-duplicate-rich than the full test set (0.406 against 0.380 at Jaccard ≥ 0.9), which if anything understates the difficulty effect. *human_nontata_promoters* sits between the regimes at $0.0345 \rightarrow 0.0898$. The linear SVM rises from rank 5 to 1 and logistic regression from 6 to 2, while 1-nearest-neighbour falls from rank 2 to 9, last, at 0.537—3.7 points above the 0.500 base rate of this balanced test set, on $n = 30,971$. The two models exchanging the top slot each lead on their own split by a paired margin excluding zero ([0.071, 0.080] as shipped, [0.024, 0.035] corrected). The four models whose drops exclude zero are 1-NN, the forest, extremely randomized trees and the MLP; the other five—both linear models, boosting, naive Bayes and 15-nearest-neighbour, which we designated memorization-prone in advance—are null. The MLP is not one of those four

Table 4. Three-arm decomposition of the corrected-split drop. **As-shipped** and **Novel-only** are the *same fitted model*—one fit on the as-shipped training set—scored on the whole as-shipped test set and on its novel stratum alone (maximum 8-mer Jaccard to any training sequence < 0.5). Their difference is **evaluation-side inflation**: it requires neither re-splitting nor retraining, and is what the benchmark overstates about a model that already exists. **Corrected** refits on the near-duplicate-aware split, and its difference from Novel-only is the **split-side effect**, a retraining penalty rather than a measurement error. The two sum to the total drop. Only the random forest pays a split-side penalty on *human.enhancers.ensembl*; the other three improve slightly when refit on de-leaked data. Corrected accuracies are means over five re-split seeds. The novel stratum's class balance differs from the full test set's (0.197 against 0.500 here), so these raw figures are prevalence-confounded and the balanced-accuracy versions are given in the text.

Dataset	Model	As-shipped (full test)	Novel-only (same fit)	Corrected (re-split)	Eval-side inflation	Split-side effect	Total drop
human.enhancers.ensembl	RandomForest	0.8596	0.7698	0.6957	+0.0898	+0.0741	+0.1639
	LinearSVC	0.7980	0.7820	0.7952	+0.0160	−0.0132	+0.0028
	LogisticReg.	0.7809	0.7649	0.7776	+0.0161	−0.0127	+0.0033
	HistGB	0.7636	0.7488	0.7574	+0.0148	−0.0086	+0.0063
human.nontata.promoters	RandomForest	0.9284	0.8793	0.8129	+0.0491	+0.0664	+0.1155
	LinearSVC	0.8876	0.8546	0.8063	+0.0331	+0.0482	+0.0813
	LogisticReg.	0.8702	0.8340	0.8300	+0.0361	+0.0041	+0.0402
	HistGB	0.8905	0.8383	0.8244	+0.0523	+0.0139	+0.0661

Model ranking: original (leaky) vs near-duplicate-aware split ($k=6$, accuracy)




Figure 3 Four-model rankings on the original (leaky) split versus the near-duplicate-aware split ($k=6$, accuracy). The random forest is dethroned on *human.enhancers.ensembl* (middle) and, under accuracy only, on *human.nontata.promoters* (left), while rankings are stable on the clean control *demo.human.or.worm* (right, 20k subsample). Kendall τ is deliberately not annotated: at four models it is too coarse to carry (§3.6), and it orders these panels opposite to the conclusions they support.

designated learners, so its presence says the inflation reaches a neural learner too.

The convolutional network.

The leaky split ranks it fifth of ten; the corrected split ranks it first. The registered nine are unaffected by adding it—the linear SVM remains the corrected winner of the registered nine—and the CNN is a tenth, *post hoc* entrant. Replicating the pre-registered reference cell across five training seeds on the same corrected partition gives corrected accuracies 0.8028–0.8131 (mean 0.8086, sd 0.0040) against the linear SVM's 0.7975, and as-shipped accuracies 0.8119–0.8359, between the forest's

0.8596 and the linear SVM's 0.7980, so the rank is the same on every seed. The exchanging pair is the MLP and the CNN, and both arms of the inversion criterion clear zero: as-shipped MLP–CNN +0.0500 [0.0320, 0.0679], corrected CNN–MLP +0.0402 [0.0309, 0.0495]. Against the harder comparator the criterion does not require, *LinearSVC*, the paired difference is +0.0111 with a combined-source interval of [0.0019, 0.0204], folding the seed-to-seed standard deviation (0.0040) into the cluster-bootstrap standard deviation (0.0026) in quadrature, because the quantity of interest is the advantage of *one* trained network rather than of an average over five. Bootstrapping seed-averaged per-example correctness instead gives [0.0063, 0.0161],

narrower than any individual seed's interval because averaging removes the training variance rather than estimating it, and we do not quote that as the result. Four of the five individual seeds have paired intervals excluding zero and the fifth does not ($+0.0053$, $[-0.0006, 0.0109]$), so the advantage is uniform in rank but not in significance. The reference cell is the *maximum* of the nine-cell regularization grid, whose corrected accuracies span 0.749 – 0.813 with 0.7975 inside—five cells above, four below—so no conclusion is drawn from that cell alone; the claim rests on seed replication of one pre-specified cell against a paired interval.

human.nontata.promoters: a dead heat.

On that dataset the picture is weaker. Under accuracy the forest is demoted from rank 1 to 3 (corrected LR $0.829 > \text{HGB } 0.827 > \text{RF } 0.820 > \text{LinearSVC } 0.808$, all seed-0; the forest's 0.820 is the seed-0 value of the quantity whose five-seed mean, 0.810 , gives the headline drop). This is a statistical dead heat—corrected CIs overlap (RF $[0.812, 0.828]$, LR $[0.821, 0.837]$)—and does not survive a threshold-free metric: under AUROC and F_1 the forest remains rank 1 (corrected AUROC 0.929 , highest of the four). With only four models Kendall τ is too coarse for a primary statistic. The nine-model roster produces a larger scramble there with supported swap margins, but the prevalence-corrected graded gap does not separate the model families on it (memorizers $+0.122$ against $+0.179$), so the wider roster does not strengthen that dataset's case; an alternative discriminator fails too, the residual of each drop against φg being larger on *human.enhancers.ensembl* ($+0.257$) than on *human.nontata.promoters* ($+0.152$), so it measures the retraining penalty rather than dataset quality. The reordering therefore stands as an existence proof on *human.enhancers.ensembl* and as a cautionary partial case on *human.nontata.promoters*.

The six registered predictions, scored. P1 fails as registered: on the raw gap the largest value belongs to the forest on *human.enhancers.ensembl* (0.230 against 1-NN's 0.208 and extremely randomized trees' 0.210) and to 15-nearest-neighbour on *human.nontata.promoters*. Under the prevalence-corrected gap 1-NN does lead on the first ($+0.461$ against $+0.319$), but that statistic was defined after P1 was committed and the reading is *post hoc*, not a confirmation. P2 holds on *human.enhancers.ensembl* (1-NN has the largest drop, 0.336) and fails on *human.nontata.promoters*, where both P1 and P2 fail. P3 holds decisively on the first (mean rank 4.50 against 5.50 as shipped, 7.25 against 1.50 corrected). P4 holds: on the clean control no model's drop interval excludes zero, for any of the nine—that control runs on a 20,000-sequence subsample of *human.enhancers.cohn* rather than its full 27,791, admissible because its full-scale leak fractions are 0.001 and 0.005 . P5 and P6 are scored in §3.10 and Supplementary §S2.

3.7. The mechanism (exploratory), and its scope

On *human.enhancers.ensembl* the near-default forest scores 1.000 on near-duplicate test sequences (≥ 0.9 similarity) but 0.770 on novel ones—a raw graded gap of $+0.230$ against 0.034 – 0.037 for the other three learners of the four-model comparison, though across the wider roster the raw gaps reach

0.210 and 0.208 . On novel sequences the forest (0.770) is in fact worse than the linear SVM (0.782). Those are the registered quantities, and on them P1 fails.

Exploratory: the prevalence-corrected gap.

Everything from here to the end of this subsection is exploratory and is reported as such. P1 was committed against the raw graded gap, and the prevalence-corrected form below was defined after that prediction had been scored, so it cannot confirm the hypothesis it was built to rescue—it can only motivate the prospective test we register at the end of this paragraph. Correcting the class-composition confound *strengthens* the mechanism: from balanced accuracy, 1-nearest-neighbour, the definitional memorizer, has by far the largest gap ($+0.461$), then extremely randomized trees ($+0.370$) and the forest ($+0.319$), with the linear models ($+0.154$, $+0.144$), naive Bayes ($+0.129$) and, lowest, 15-nearest-neighbour ($+0.016$)—averaging over fifteen neighbours is what stops a nearest-neighbour rule memorizing. The four memorization-prone learners average $+0.292$ against $+0.165$, a separation the raw statistic does not produce. Memorization pays because near-duplicates carry the label: at ≥ 0.9 similarity 99.99% of near-duplicate test sequences share their nearest training neighbour's label on *human.enhancers.ensembl* ($n = 11,774$) and 99.95% on *human.nontata.promoters* ($n = 2,037$); clean datasets have near-empty high-similarity bins (Supplementary Fig. S2). The correction is equally informative where it removes a signal: on *human.nontata.promoters* the corrected gaps do *not* separate the families—memorizers average $+0.122$ against $+0.179$ for the other five, the largest belonging to the MLP—so the mechanism claim is scoped to *human.enhancers.ensembl*. We therefore register the corrected gap forward rather than counting it backward: on the next leaky dataset audited, 1-nearest-neighbour will carry the largest prevalence-corrected graded gap and the four memorization-prone learners will average a larger corrected gap than the remaining five. That prediction is falsifiable on one dataset, and *human.nontata.promoters* above is the case in which it would already have failed.

Exploratory: the shape of the curve, against a published one.

That accuracy grades with training-set similarity, and that models rely on memorized associations where it does, is established by Rafi et al. (2026) across three modelling regimes; it is not claimed as novel here, and what this suite adds is a matched clean control and an editable construction step. Their curve is non-monotonic—best at both extremes of homology, worse at intermediate similarity—so ours is worth reporting whether or not it agrees. Binning test sequences by maximum similarity to training and reading *balanced* accuracy, so the prevalence confound above cannot manufacture a shape, the two suites disagree in an interpretable way. The memorizing forest is monotone increasing on both leaky datasets ($0.681 \rightarrow 0.818 \rightarrow 0.910 \rightarrow 1.000$ on *human.enhancers.ensembl*; $0.788 \rightarrow 0.922 \rightarrow 0.997 \rightarrow 1.000$ on *human.nontata.promoters*), which is what a memorization account predicts and what their non-monotonic curve does not show. The models that do *not* memorize are non-monotonic on both, but in opposite directions: on *human.nontata.promoters* logistic regression dips at intermediate similarity ($0.824 \rightarrow 0.755 \rightarrow 0.800 \rightarrow 0.963$), the shape they report, while on

human_enhancers_ensembl it peaks there instead (0.747 $\rightarrow$ 0.955 $\rightarrow$ 0.886 $\rightarrow$ 0.901). Two things limit this. The intermediate bins are small on *human_enhancers_ensembl*—311 and 115 sequences against 18,770 and 11,774 at the extremes—so the peak rests on few observations; and their regimes reach 196 kb contexts and functional divergence between homologs, neither of which a k -mer model on ≤ 573 bp windows can express. We report the difference rather than resolve it (`results/graded_shape_summary.csv`).

3.8. The effect is a property of the unregularized forest

At `min_samples_leaf=1` the forest scores exactly 1.000 on the ≥ 0.9 bin, with drops of 0.164 and 0.108 (seed 0), graded gaps 0.230 and 0.121, and rank 1. At 500 the drop is ≈ 0 (0.164 $\rightarrow$ -0.001; 0.108 $\rightarrow$ 0.000) and the leaf-size gap collapses (0.230 $\rightarrow$ 0.020; 0.121 $\rightarrow$ 0.061), monotonically but for one 0.0012 reversal in the last *human_nontata_promoters* step; along the depth axis the gap is non-monotone there, rising to 0.224 at depth 16 before falling to 0.161. Crucially, the “forest wins on the leaky split” phenomenon itself *requires* the unregularized forest: as shipped it holds rank 1 only near the default and has fallen to rank 4 by `min_samples_leaf` ≥ 20 on both datasets. Regularization removes the inflation without recovering the accuracy the inflation concealed—corrected accuracy on *human_nontata_promoters* falls 0.820 $\rightarrow$ 0.785 along the same path—so the two interventions are not interchangeable.

3.9. Deep models, binary and multiclass

A from-scratch 1D CNN trained only on each training split shows the same inflation, so the effect is not a classical-model artifact. The pre-registered reference cell drops on *both* datasets—+0.043 on *human_nontata_promoters* [0.025, 0.061], +0.014 on *human_enhancers_ensembl* [0.007, 0.021]—so the refutation condition, which requires the interval to include zero on *both*, is not triggered. One *human_enhancers_ensembl* seed does include zero, and *human_nontata_promoters*, which carries the larger drop, was not seed-replicated and is the more dispersed of the two across the grid, so the verdict rests on an unreplicated cell on the dataset from which §3.7 scopes the mechanism away. Nineteen of 20 cells drop with intervals excluding zero; the exception is the *human_nontata_promoters* dropout-0.6 cell, where accuracy *rises* 0.020. Seed replication of the *human_enhancers_ensembl* reference cell gives drops of 0.0049–0.0269, a five-fold spread whose smallest does not exclude zero, so +0.014 is one draw and no single-run drop here, including each of the nineteen, is a point estimate. One registered expectation was not met: the pre-registration named this a *dose-response* grid, and it is not one, the drop being non-monotone in both dropout and weight decay.

The multiclass question is answered by construction, and the two model families answer it differently. On the identical injected construction the forest’s drop grows monotonically with f : +0.0004, +0.1177, +0.1791, +0.2581 at $f = 0, 0.1, 0.2, 0.4$. Logistic regression is not untouched—it drops 0.0158, 0.0323, 0.0661, monotone in f , its as-shipped accuracy climbing 0.591 $\rightarrow$ 0.669 as it exploits the injected copies—but the

forest’s advantage is $7.4\times$ to $3.9\times$ larger, narrowing as the dose rises. The CNN shows *no* dose-response: mean drops -0.0386 (sd 0.0239), -0.0182 (0.0296), -0.0249 (0.0325), -0.0089 (0.0362) over three seeds at the same four doses (each cell a single run)—every mean negative, the seed standard deviation equalling or exceeding any dose effect. All three seeds are reported; seed 1 underfits, its as-shipped accuracy averaging 0.6856 against 0.7552 and 0.7534, and is retained rather than excluded. The mechanism therefore generalizes to multiclass *for the classical memorizer* and not for the regularized network—what “memorization propensity, not capacity” predicts, since the label-carrying copies are available to every learner and the differential is a property of which learner exploits them.

3.10. Ordinary tuning selects the memorizing configuration

A practitioner does not know a split is leaky; they tune by cross-validating on the training set. On *human_enhancers_ensembl*, ordinary stratified 5-fold selects `min_samples_leaf=1`, the memorizing default, and the model it selects scores 0.860 as shipped but 0.696 corrected—a loss of 16.4 points. Cluster-grouped 5-fold selects 20, and that model loses only 2.2 points (0.751 $\rightarrow$ 0.729). Worse, the benchmark punishes the correct methodology: the leakage-aware tuner’s model appears 10.9 points *worse* as shipped (0.751 against 0.860) while being 3.3 points *better* corrected (0.729 against 0.696)—a point-estimate comparison to which we attach no interval. On *human_nontata_promoters* both procedures select 1, so P5 holds on the first dataset and is a tie, not a reversal, on the second. The untuned default is therefore not a careless choice a tuner would correct; it is what ordinary model selection selects, and the benchmark rewards it.

3.11. A second suite carries the same defect: three of eleven independent tasks are leaky

Counting independent tasks, three of eleven Nucleotide Transformer tasks are leaky and six more borderline, leaving two clean. *enhancers* carries 25.0% test-to-train leakage that is *byte-identical*—verified on a code path using no k -mers, 100 of 400 test sequences at label concordance 1.000—while *splice_sites_acceptors* and *donors* carry 0.244 and 0.245 at Jaccard 0.7 with zero exact duplicates. An earlier pass capped training sets at 20,000 and returned three borderline and five clean; re-running uncapped moved three histone tasks across the cut on containment.

The splice mechanism is not the cross-species homology an earlier account proposed: 65.9% and 68.3% of leaked test sequences share the same UniProt entry name with their best training partner—same gene, same organism—and at Jaccard ≥ 0.9 that share *rises* to 77.1% and 82.6%, where homology predicts thinning. It is same-gene splice-window redundancy, and label concordance among leaked pairs is correspondingly poor (0.819, 0.792 against 0.999 on both Genomic Benchmarks datasets). That account is offered after the fact, not predicted (Supplementary §S6).

Running the full pipeline there returns a null: no material ranking inversion on any leaky task nor on the control. Leakage does inflate scores on *enhancers*—the forest drops 0.0413 [0.0241, 0.0583], *LinearSVC* 0.0221 [0.0055, 0.0381], *HistGradientBoosting* 0.0185 [0.0021, 0.0366], three of the four learners excluding zero, while logistic regression, leading on both splits, is null (+0.0046 [−0.0104, 0.0199]). The memorization *ordering* reproduces—the forest carries the largest drop and is the only learner with a positive graded gap (+0.100 against −0.084 to −0.169)—but the confinement to memorizers reported on *human.enhancers.ensembl* does not: two non-memorizers exclude zero here, and their negative graded gaps rule out memorization, leaving the retraining penalty of §3.6 the better reading. On *splice_sites.acceptors* the memorizing forest's drop is null (+0.0201 [−0.0068, 0.0446]), so the model carrying the mechanism is the one the effect does not reach. On *splice_sites.donors* correction *raises* accuracy for three of four learners and one gain is nominally supported (*LinearSVC* −0.0251 [−0.0526, −0.0004]), marginally and seed-dependently ($p \approx 0.059$ uncorrected). Under Benjamini–Hochberg across this cross-suite grid only *enhancers*|RF survives, so these flags are nominal—with one qualification that runs against us and is therefore reported: under a single pooled family of all 31 deltas (§3.13) *enhancers*|*LinearSVC* survives as well, at $q = 0.049$, close enough to the cut that the normal approximation behind these p -values cannot resolve it either way. The leading model is unchanged on all four tasks, which is the claim this section rests on: reordering additionally requires a challenger genuinely better on novel sequences, and this suite supplies leakage without one.

3.12. A pre-registered cross-suite test, falsified

Our pre-registration carried binding predictions for GUE (Zhou et al., 2023b), recorded before that suite was examined—the prediction document was committed as 2edad5c on 2026-07-18 and the census that scored it as 3af3eb4 nine hours later, in separate commits whose ordering is published with the `git log` output in `PREREGISTRATION.md`, subject to the limits of author-controlled timestamps stated there: that its short fixed-length human regulatory tasks—core promoter at 70 bp, promoter at 300 bp, transcription-factor binding at 101 bp—would be leaky, and its multi-species tasks clean. At full scale the registered predictions score 3 of 15, and all eleven predicted-leaky tasks are clean, measuring 0.009 to 0.041 under Jaccard and 0.014 to 0.042 under containment—little leakage, not none, the largest under half the verdict cut. These are not eleven independent benchmarks: the *_all* promoter tasks' test sets are the exact union of their own *notata* and *tata* test sets, and the 70 and 300 bp promoter families are the same loci at two window widths, leaving seven independent test partitions, which bounds how much independent evidence the null carries.

This refutes the hypothesis the prediction encoded. Leakage does not track task type: seven independent benchmarks in precisely the regime we flagged do not carry it. What it tracks is the construction step of §3.2—Genomic Benchmarks' two leaky datasets have an un-deduplicated assembly where GUE's comparable tasks do not. Two verdicts moved. *emp.H3K4me3*, registered clean, has a full-scale containment leak fraction of

0.138 against 0.093 capped, making it borderline: under the binary rule committed in advance it remains a correct prediction and the score stands at 3 of 15; under a three-way rule it is a miss and the score is 2 of 15. And *virus.covid*, also registered clean, measures 1.000—the twelfth miss, and the one discussed in §4, where it is not a curation defect. Overall, the prediction that leakage follows task type is refuted by seven independent partitions in precisely the regime it named; the two verdicts that moved are borderline calls, not reversals.

3.13. Robustness

The block bootstrap.

That the intervals move little is not because the dependence is absent. The splitter assigns whole components to one side, so every near-duplicate relation surviving correction is test-to-test by construction, and the corrected test sets are the more clustered arm: 43.4% of corrected *human.nontata.promoters* test sequences lie in a multi-member component against 25.2% as shipped, and 47.7% of *human.enhancers.ensembl*'s against 10.1%. Those sizes are heavy-tailed, so Kish's m^* is 4.45 and 2.58 where the mean is 1.55 and 1.33, giving design effects of 2.80 and 2.58 for the forest where the mean-size convention returns 1.29 and 1.33. The inflation is measured directly: on the corrected arm the cluster-bootstrap variance of the forest's accuracy exceeds its sample-bootstrap variance by $3.96\times$ and $2.12\times$ ($3.10\times$ and $1.78\times$ for logistic regression), widening the delta intervals $1.2\text{--}1.8\times$. The intraclass correlation is high (0.78–0.99) across the four fits, so the design effect is governed by the block-size distribution rather than by the within-block correlation. The dependence is therefore real, larger than a mean-size design effect suggests, concentrated in the arm the correction creates—and already inside every reported interval, because the block bootstrap resamples those components rather than modelling them. No interval and no verdict changes: the forest's drop still excludes zero on both datasets ([0.155, 0.173]; [0.092, 0.126]), and Benjamini–Hochberg over the declared 15-delta family leaves the same set significant. One source no test-set resampling reaches is the fit, which is held fixed.

Chromosome holdout.

The demotion reproduces on *human.enhancers.ensembl*: the forest falls from rank 1 (0.8596) to rank 4 (0.7069), giving the identical corrected order as the near-duplicate-aware split. On *human.nontata.promoters* it does *not*: the forest stays rank 1 (0.9284 $\rightarrow$ 0.8199), a third independent signal—with the AUROC/ F_1 reversal and the novel-only rank correlation—that its demotion is accuracy-only. Two disclosures bound that arm: its test set is 74.33% positive against a 47.20% training prevalence, a 27-point shift, and it is scored on accuracy, which the imbalance panel below shows understates the correction. The controls agree on the drop magnitude on *human.enhancers.ensembl* (corrected 0.707 against 0.696); on *human.nontata.promoters* they agree to 0.0001 across test sets 74.3% and 54.4% positive, which is consistent with, but no evidence for, the near-duplicate-aware split proxying the deployment control.

Metric geometry and cluster structure.

The demotion survives canonical (reverse-complement-collapsed) k -mer features (forest 1 $\rightarrow$ 4, drop 0.153

on *human_enhancers_ensembl*; $1 \rightarrow 3$, drop 0.097 on *human_nontata_promoters*), so it is not forward-strand overfitting, and a GC-in-distribution restriction gives lower corrected accuracy (0.787, 0.631), not the inflated original, so it is not a GC covariate shift. An MMseqs2 alignment-scored re-clustering reproduces the effect on *human_enhancers_ensembl* (0.164 against 0.162) and gives a milder drop on *human_nontata_promoters* (0.108 $\rightarrow$ 0.064); that arm changes clustering engine and linkage rule together, so the difference is a linkage signature as much as a metric one (Supplementary §S3). Cluster cohesion explains the asymmetry between the two datasets, and geometrically rather than biologically: *human_enhancers_ensembl*'s clusters are 99.6% pairwise at cohesion 0.96, discrete $2\times$ locus duplication that whole-cluster correction removes cleanly, whereas *human_nontata_promoters*'s 420-member largest component sits at cohesion 0.083, is 420/420 negative class, and is a single locus tiled by 251 bp windows at a median 2 bp step—window geometry, not chained paralogy. That geometry has a documented origin rather than an inferred one, though not the one a reader might guess. The negatives are read from a third-party FASTA that already contains 27,731 uniformly 251 bp records, and the curators' seq2loc step only recovers their GRCh38 coordinates by exact string match; it does not create the windows and nothing merges them afterwards. The tiling is visible in the notebook's own printed output, whose consecutive chr1 negatives sit at steps of 23, 8, 41 and 32 bp—eight rows of a preview, so corroboration of our census rather than a census in itself. Dataset-wide 94.3% of its negatives sit in multi-member components against 0.4% of positives and 2,808 of 2,811 components are class-pure, the same fact §3.2 reports as $R_{\text{self}} = 0.961$ in the negative class; the certification tool therefore returns *leaky-but-correction-unsafe* rather than *leaky* below cohesion 0.5. A dustmasker check excludes a microsatellite artifact but is silent on interspersed repeats, and on *human_nontata_promoters* the leaked stratum is in fact $\sim 7\times$ AluY-enriched relative to novel (Supplementary §S3).

Where the multiplicity family is drawn.

Correcting the per-dataset grid and the cross-suite grid as separate families is more permissive than correcting their union, so we ran the union. Pooling all 31 deltas into one Benjamini–Hochberg family leaves every conclusion in this paper unchanged and adds one flag rather than removing any: *enhancers*|LinearSVC crosses at $q = 0.049$. Pooling is not the one-sided stress test it looks like—enlarging a family raises the number of tests, which is restrictive, but it also raises the rank of each p -value when the incoming tests are more significant than the incumbents, and the per-dataset family contributes several p -values below 10^{-5} . Both effects are real and the analysis has to be run rather than reasoned about (*results/bh_pooled.csv*).

Against a published tool.

hashFrag (Rafi et al., 2026) is the closest runnable comparator: BLAST-based, alignment-scored, and strand-explicit, building its database from both orientations of every training sequence. Run on all eight datasets, it agrees on the two verdicts that matter and disagrees in an instructive way about the rest. Its threshold is the whole story. hashFrag supplies no

default and its documentation recommends scoring dinucleotide-shuffled sequences and cutting above that null; done per dataset, that procedure returns thresholds of 21–88 and calls every dataset leaky, including all six we call clean. That is the failure our own supplement predicts on independent grounds—real genomic DNA carries repeats and low-complexity tracts that shuffling destroys, so a synthetic null puts the cut far below the scale of real alignment scores—and it is a caution about the recommended procedure rather than about the tool. Sweeping the threshold instead, a window at 220–240 reproduces 7 of our 8 verdicts exactly, with per-sequence agreement of 0.837 on *human_enhancers_ensembl* and 0.958 on *human_nontata_promoters* and Cohen's κ of 0.636 and 0.913 respectively. The separation is in the right direction and is not an artifact of choosing that window: the two leaky datasets require thresholds of 260 and 380 before their leak fractions fall under the verdict cut, against 100–220 for the clean ones.

The single persistent disagreement is *drosophila_enhancers_stark*, which hashFrag calls leaky at every threshold we swept while we call it clean, and it marks the real boundary of our instrument rather than an error in either tool. Its sequences are long (2,142 bp median) and tile the genome, so neighbouring tiles share flanking sequence over hundreds of bases: BLAST scores those alignments highly, while an 8-mer Jaccard is bounded by the length ratio and stays low. The dataset is clean of near-duplication and is not clean of homology—the same distinction the measured recall of §4 quantifies, arrived at through a second tool and a different algorithm. Symmetrically, a single absolute alignment score is itself length-dependent: at threshold 220 hashFrag reads lower than we do on *human_enhancers_ensembl* (0.271 against 0.384), because that dataset's short sequences cannot reach 220 at all. Neither instrument is threshold-free, and both should have their cut follow sequence length.

Imbalance, and what “corrected” means.

Under induced prevalence imbalance the inflation is revealed by prevalence-aware metrics and masked by accuracy: on *human_enhancers_ensembl* at $\pi = 0.2$ correction lowers AUPRC 0.622 $\rightarrow$ 0.477, MCC 0.422 $\rightarrow$ 0.182 and minority recall 0.258 $\rightarrow$ 0.070 while accuracy barely moves (0.844 $\rightarrow$ 0.808), and on *human_nontata_promoters* AUPRC falls 0.949 $\rightarrow$ 0.813 and MCC 0.785 $\rightarrow$ 0.682 against accuracy's 0.934 $\rightarrow$ 0.903—but minority recall is flat there (0.673 $\rightarrow$ 0.674), so the recall collapse does not replicate. What generalizes is that AUPRC and MCC fall several times further than accuracy on both datasets, so accuracy at threshold 0.5 understates the correction; the panel is random-forest-only and its prevalence induced rather than native, leaving the reordering claim untested under imbalance. Across the four robustness arms—block bootstrap, chromosome holdout, metric geometry and induced imbalance—no interval and no verdict changes, and the dependence the block bootstrap measures is real, larger than a mean-size design effect suggests, and already inside every reported interval.

4. Discussion

Near-duplicate leakage in this suite is produced by a step in how the data were assembled, and that step is visible, editable and reversible before any model is trained: a coordinate signature carrying no sampling uncertainty and no model dependence,

a manipulation that removes and imposes the defect at will, and a pre-registered task-type prediction that failed on all eleven targets. On one dataset the downstream cost is that the benchmark reports the wrong winner.

The reordering requires the *unregularized* forest, which by `min_samples_leaf` ≥ 20 has already fallen to rank 4 as shipped (§3.8). That configuration is not a careless choice a practitioner would correct: ordinary stratified cross-validation on the as-shipped training set selects it, and the benchmark then penalises the practitioner who tunes correctly, showing the leakage-aware tuner's model 10.9 points worse as shipped and 3.3 points better corrected (§3.10). The property that matters in a benchmark is anyway not whether some user tunes but whether the data reward memorization at all, and §3.3 locates that reward in the construction step. The controlled variable is memorization propensity, not capacity: the nominally higher-capacity HistGradientBoosting is barely inflated where the near-default forest is heavily inflated, and the regularized CNN shows no dose-response under injection at all—its mean drop is negative at every dose—where the near-default forest's grows monotonically to +0.258 (§3.9).

The suite's own published baseline is a convolutional network (Grešová et al., 2023), not a tree ensemble. A CNN-based leaderboard is affected in the direction that costs the architecture: our from-scratch network ranks fifth of ten as shipped and first corrected, scoped to one dataset and to a *post hoc* comparison, with one of its five seeds not individually clearing zero (§3.6). Its inflation is between six- and thirtyfold smaller than the forest's depending on seed, so the ratio is seed-dependent and is reported as a range; the reordering condition is not met by a CNN-only comparison, and the result is bounded by the comparison set: ten learners, one dataset.

One interpretive point governs every corrected number in this paper. A near-duplicate-aware split answers how a model scores on test sequences with no near-copy in training, which is not how a deployed model will perform, since paralogues, gene families and conserved regulatory elements are real inputs it will meet, and whole-cluster assignment also strips near-duplicate partners from the *training* set (§3.6). "Corrected" therefore qualifies the split and not the accuracy: a corrected accuracy is a lower bound on deployment performance under one adversarial notion of novelty. The as-shipped ranking is wrong; the corrected accuracy is a bound, not a deployment estimate.

human_enhancers_ensembl is an existence proof: the ranking change holds under accuracy, AUROC, F_1 , a bootstrap winner probability, a chromosome-holdout control and an alignment-scored re-clustering. *human_nontata_promoters* is a cautionary partial case, five independent signals showing its demotion is accuracy-only: overlapping corrected CIs, the AUROC/ F_1 reversal, the chromosome control, the novel-only rank correlation and the milder alignment drop. Its clusters are the geometric case rather than the duplication case, so whole-cluster correction is a coarser instrument there. Because the effect is dataset-specific the contribution is constructive: not that these benchmarks are broken, but which datasets to trust and a tool to certify a new one. Green verdicts carry a heavier evidential burden than red ones, and Table 1 states what a "clean" entry does and does not certify.

Scale matters practically. Subsampling can mask leakage—*human_enhancers_ensembl* appears almost clean at 20k

sequences and is heavily leaky at full scale—so audits must run at full dataset scale. Both censuses reported here were capped at 20,000 training sequences on a first pass and were wrong the same way: re-run in full, GUE's leak fractions rose by up to $2.8\times$ on the tasks large enough to have been capped and one verdict moved clean to borderline, and three Nucleotide Transformer histone tasks crossed the cut, leaving three of eleven independent tasks leaky and six borderline (§3.11, §3.12).

The evidence spans more than regulatory-element classification (Fig. 4): the fifty-one dataset-task censuses cover six task types and six source organisms, and forty-four of the fifty-one fall left of the verdict cut, including all eleven GUE tasks registered in advance as leaky. Length is the weaker axis: no censused task exceeds a 1,000 bp median in sequence space, and the one long-range datapoint is in coordinate space only (§3.4). Protein sequence is out of scope for this detector: an 8-mer over a twenty-letter alphabet collapses the Jaccard towards exact match, so a protein census would need k re-chosen and the calibration redone—a different instrument.

What the detector detects, and one false positive.

The 8-mer Jaccard at 0.7 is not a fixed instrument across length (§2.4), so every leak fraction is a lower bound whose tightness depends on length. Specificity is measured rather than argued. Adjudicating a band-stratified sample of test–train pairs by Smith–Waterman against the CD-HIT-EST-like criterion the curators omitted ($\geq 50\%$ coverage of the shorter sequence at $\geq 80\%$ identity), every flagged band on both leaky datasets is 100% homologous: precision 1.000 (Clopper–Pearson lower bounds 0.951 and 0.978) and a per-pair false-positive rate of 0.000, bounded above by 1.5×10^{-7} and 1.1×10^{-6} —well below the $1/n_{\text{train}}$ scale at which a maximum-over-training-set statistic would be corrupted. Recall runs the other way, 0.006 and 0.260, because homology at 80% identity over half a sequence routinely leaves the 8-mer Jaccard below 0.7; that is the quantitative form of the lower-bound scope above, not a defect (Supplementary §S4). Construction and provenance now corroborate the measurement rather than substitute for it: the merged arm of fix-a-leaky is clean as a property of the edit itself and reads 0.012, the dose-zero control 0.004, and the five clean Genomic Benchmarks datasets, whose *upstream resources* are merged—the Ensembl Regulatory Build for *human_ocr_ensembl* and Cohn et al. (2018) for *human_enhancers_cohn*, not any step in Genomic Benchmarks' own construction, which documents none (§3.2)—0.001 to 0.016, against the two leaky datasets' 0.384 and 0.406. The GUE tasks cannot join that list: their "clean" label is this estimator's output, so quoting it would be circular. What they supply independently is the exact-duplicate rate, by string equality with no k -mers, running 0.000 to 0.026 across the eleven and agreeing with the k -mer verdicts through a separate code path.

The detector also cannot distinguish a curation defect from biology. GUE's *virus_covid*, registered clean, measures a near-duplicate fraction of 1.000, the largest we observed anywhere, and an exact-duplicate rate of 0.122—yet it is not a curation defect. The task is nine-way SARS-CoV-2 variant classification over 999 bp windows of a ~ 30 kb genome whose variants differ by a handful of mutations, so test-to-train similarity has a *minimum* of 0.824 and a median of 0.987: no split could be otherwise, and a leak fraction carries no information




Figure 4 Near-duplicate leak fraction across the fifty-one dataset–task censuses reported here, in three benchmark suites. Each point is one censused dataset–task row, which is not the same as an independent task, and the panel annotates both conventions: the Nucleotide Transformer contributes 13 shipped tasks in two releases (26 rows) that collapse to 11 independent tasks, GUE ships 17 tasks of which the 11 registered in advance collapse to 7 independent test partitions, and the text reports independent-task counts throughout. x is the fraction of test sequences whose maximum 8-mer Jaccard to any training sequence exceeds 0.7, measured at full scale; position on the axis is the *square root* of that fraction, as the axis label states, so the sub-0.1 region, which holds forty-four of the fifty-one, is legible. Marker shape denotes task type (13 promoter, 7 enhancer, 6 splice site, 13 histone-mark or nucleosome occupancy, 7 transcription-factor binding, and five singletons: open chromatin, coding-versus-intergenomic, human-versus-worm, three-class regulatory assignment and nine-way viral variant assignment); colour denotes source organism. The dashed rule at 0.1 is the LEAKY verdict cut used throughout. Open symbols are the eleven GUE tasks registered in advance as leaky; all eleven fall left of the cut, at 0.009 to 0.041. The single grey-haloed point is *human_nontata_promoters*, the one dataset certified *leaky-but-correction-unsafe* (Table 1). The SARS-CoV-2 point at 1.000 is *virus_covid*, where near-identity is a property of the genome rather than a curation defect (§4). A fourth suite, BEND, is censused in coordinate space only and appears in Fig. 1B.

there. Any application to viral, organellar or highly conserved sets must establish the distinction by other means—and the coordinate remedy is unavailable on viral releases, which ship no coordinates.

Pretrained models and the second leakage channel.

Those models are contaminated by construction, so fine-tuning them here would produce a confounded number rather than evidence: the assembly is verified GRCh38 (§2.2). The overlap is nonetheless not total, and the exception is useful. HyenaDNA inherits the Enformer interval set (Avsec et al., 2021) and designates chromosomes 14 and X as its held-out pretraining test chromosomes; both are present here, carrying 6.6% of *human_enhancers_ensembl*'s test set and 15.3% of *human_nontata_promoters*'s, the latter unevenly by class (24.6% of test negatives against 7.5% of positives). Pretraining overlap is therefore at most ~93.4% and ~84.7% for that model, not ~100%; for DNABERT-2 and the Nucleotide Transformer, whose human component is one part of a multispecies corpus, the argument remains one from provenance.

One logical point cuts in our favour: pretraining is label-free and common to both split arms, so it *attenuates* the as-shipped-minus-corrected difference rather than inventing it, and a positive finding under contamination is conservative evidence. What would settle the question is an overlap-free arm, and chromosomes 14 and X are a stratum HyenaDNA provably never saw—the right future protocol, and one these datasets already contain enough of to carry an interval. A

second route, priced but not performed, applies break-a-clean to *drosophila_enhancers_stark* and fine-tunes HyenaDNA on both arms under both splits, manufacturing the contamination inside a genome the model has never seen: 6,914 sequences of at most 3,237bp and twelve runs of *hyenadna-small-32k*. Whether deployed foundation models inherit this inflation on human benchmarks is left open; our from-scratch CNN establishes only that the mechanism reaches a trained neural network with no pretraining confound.

The obvious next targets—DeepSTARR (de Almeida et al., 2022), BEND (Marin et al., 2024) and DART-Eval (Patel et al., 2024)—sharpen our prediction rather than confirm it. Each assembles its examples from genome-tiled or annotation-overlapping intervals, the construction step that produced the defect here, but each also partitions by genomic position or sequence identity: DeepSTARR bins STARR-seq peaks into 249bp windows at a 100bp stride yet holds out the two halves of chromosome 2R; DART-Eval holds out chromosomes 5, 10, 14, 18, 20 and 22; BEND partitions by chromosome on three tasks and at 80% identity on the fourth. Overlapping intervals therefore cannot span those splits, and our condition predicts all three clean—confirmed directly for BEND in §3.4. The condition our evidence supports is thus conjunctive: overlapping construction *and* a split that ignores genomic position—the pair this paper's title names, and the reason neither the construction step nor the split rule is separately diagnostic. Our own suite makes the same point, since *drosophila_enhancers_stark* tiles

at a 2,142bp median and audits clean because its cross-split overlaps are base-pair-scale touches rather than shared windows.

5. Conclusion

Near-duplicate leakage in a widely used genomic benchmark suite is a curation defect with a signature in the shipped coordinates, and it does more than inflate scores: on one dataset it changes which model the benchmark reports as best. Two of the seven binary datasets are leaky—one of them only in the weaker *correction-unsafe* sense—one borderline under a length-robust metric and four clean; the cause is 77,421 positive intervals on 40,934 distinct coordinates (§3.2), and editing that construction step alone both abolishes the syndrome and manufactures it in a clean dataset (§3.3). Correcting the split costs the best model 16.5 points on *human_enhancers_ensembl*, combined-source [15.4, 17.4], and 10.8 points on *human_nontata_promoters*, [8.8, 12.9]—each interval quoted on the estimand it was computed on, the latter the seed-0 partition whose five-seed mean is the 11.8 points §3.5 reports—both excluding zero where every clean delta includes it, and on *human_enhancers_ensembl* a nine-learner roster cannot separate its top three as shipped yet those three span 23 accuracy points corrected—against 14 under a difficulty-matched control, so the correction separates them by more than the change in operating point alone explains (§3.6). *human_nontata_promoters*, where the demotion is accuracy-only, is a cautionary partial case rather than a second demonstration.

Four practices follow. Audit at full dataset scale, since subsampling can make a heavily leaky dataset look clean—a subsample is admissible to confirm cleanliness, never to establish it. Report both a length-blind and a length-robust similarity, since containment flagged as borderline one dataset the Jaccard scored clean, and let the threshold follow sequence length. Partition in coordinate space before splitting, since the defect here was introduced upstream of any sequence-level analysis—though coordinates may escalate a verdict and never clear one, and an exact-coordinate collapse is a no-op on datasets like *human_nontata_promoters*, whose redundancy is in overlapping windows rather than repeated rows. And include at least one memorization-prone learner in any comparison, since the gap between its leaky and corrected scores is a direct tripwire for a contaminated split, costlier than the as-shipped screens because it requires building the corrected split. All four transfer unchanged to multiclass, where per-class prevalence should additionally be checked after splitting.

The reordering result is a single-dataset existence proof, scoped to comparisons containing a memorization-prone model. What generalizes is the cause: near-duplicate leakage here is produced by an un-deduplicated assembly step, is visible in shipped coordinates before any model is trained, can be removed and imposed at will with the downstream consequences following in both directions, and is absent from the eleven short human regulatory tasks—seven independent partitions—that a task-type hypothesis predicted would carry it. Audit the construction step and the split rule together, not the sequence type.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Conflict of Interest

None declared.

Responsible Disclosure

This paper reports a defect in a publicly maintained resource, so the finding is being reported to that resource's maintainers as well as published. A disclosure describing the duplicated-coordinate construction, its measured downstream effect and a suggested fix has been prepared for the Genomic Benchmarks issue tracker (https://github.com/ML-Bioinfo-CEITEC/genomic_benchmarks); its text is released with our code as DISCLOSURE_DRAFT.md. *It had not been filed at the time of writing*, and we state that rather than imply otherwise; the filing date, issue number and any response will be recorded here on submission. Four of the seven binary datasets carry none of this defect, and nothing here bears on the collection as a whole.

Data Availability

All code required to reproduce the analyses, along with the report card, figures, the near-duplicate-aware splitter and the certification tool, is available at <https://github.com/rihinkavuru/homology-leakage-genomic-benchmarks>. The Genomic Benchmarks datasets analysed here are publicly available and were obtained via the genomic-benchmarks package (Grešová et al., 2023); they are not redistributed in our repository. The coordinate-space census of BEND (§3.4) uses only the bed annotation files distributed with that resource (Marin et al., 2024), which carry their own split assignments and require no reference genome; those files are likewise not redistributed, and the script that downloads them and reproduces the census is included. Every upstream reference-genome and annotation URL on which the datasets depend—Ensembl releases 97 and 100, and the UCSC hg38/hg19 REST endpoints—was re-resolved on 2026-07-20 and returned HTTP 200; the release, artifact, live URL, archive host and access date for each are tabulated in `results/reference_provenance.csv`. Because Ensembl release 116 (June 2026) is the final release served on the current platform, we cite the dated archive hosts (<https://jul2019.archive.ensembl.org>, <https://apr2020.archive.ensembl.org>) alongside the live paths.

References

- Ž. Avsec, V. Agarwal, D. Visentin, J. R. Ledsam, A. Grabska-Barwinska, K. R. Taylor, Y. Assael, J. Jumper, P. Kohli, and D. R. Kelley. Effective gene expression prediction from sequence by integrating long-range interactions. *Nature Methods*, 18(10):1196–1203, 2021. doi: 10.1038/s41592-021-01252-x.

- B. Barz and J. Denzler. Do we train on test data? Purging CIFAR of near-duplicates. *Journal of Imaging*, 6(6):41, 2020. doi: 10.3390/jimaging6060041.
- Y. Benjamini and Y. Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995. doi: 10.1111/j.2517-6161.1995.tb02031.x.
- A. Bushuiev, R. Bushuiev, J. Sedlar, T. Pluskal, J. Damborsky, S. Mazurenko, and J. Sivic. Revealing data leakage in protein interaction benchmarks. *arXiv preprint*, 2024. doi: 10.48550/arXiv.2404.10457.
- D. Cohn, O. Zuk, and T. Kaplan. Enhancer identification using transfer and adversarial deep learning of DNA sequences. *bioRxiv preprint*, 2018. doi: 10.1101/264200.
- H. Dalla-Torre, L. Gonzalez, J. Mendoza-Revilla, N. Lopez Carranza, A. H. Grzywaczewski, F. Oteri, C. Dallago, E. Trop, B. P. de Almeida, H. Sirelkhathim, et al. The nucleotide transformer: Building and evaluating robust foundation models for human genomics. *Nature Methods*, 2024. doi: 10.1038/s41592-024-02523-z.
- B. P. de Almeida, F. Reiter, M. Pagani, and A. Stark. DeepSTARR predicts enhancer activity from DNA sequence and enables the de novo design of synthetic enhancers. *Nature Genetics*, 54(5):613–624, 2022. doi: 10.1038/s41588-022-01048-5.
- C. Deng, Y. Zhao, X. Tang, M. Gerstein, and A. Cohan. Investigating data contamination in modern benchmarks for large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 8706–8719, 2024. doi: 10.18653/v1/2024.naacl-long.482.
- A. Elangovan, J. He, and K. Verspoor. Memorization vs. generalization: Quantifying data leakage in NLP performance evaluation. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, pages 1325–1335, 2021. doi: 10.18653/v1/2021.eacl-main.113.
- K. Grešová, V. Martinek, D. Čechák, P. Šimeček, and P. Alexiou. Genomic benchmarks: a collection of datasets for genomic sequence classification. *BMC Genomic Data*, 24(1):25, 2023. doi: 10.1186/s12863-023-01123-8.
- U. Hobohm, M. Scharf, R. Schneider, and C. Sander. Selection of representative protein data sets. *Protein Science*, 1(3): 409–417, 1992. doi: 10.1002/pro.5560010313.
- A. Huang, I. S. Knight, and S. Naprienko. Data leakage and redundancy in the LIT-PCBA benchmark. *arXiv preprint*, 2025. doi: 10.48550/arXiv.2507.21404.
- S. Kapoor and A. Narayanan. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9): 100804, 2023. doi: 10.1016/j.patter.2023.100804.
- D. R. Kelley, J. Snoek, and J. L. Rinn. Basset: learning the regulatory code of the accessible genome with deep convolutional neural networks. *Genome Research*, 26(7): 990–999, 2016. doi: 10.1101/gr.200535.115.
- D. Kovtun, M. Akdel, A. Goncarenko, G. Zhou, G. Holt, D. Baugher, D. Lin, Y. Adeshina, T. Castiglione, X. Wang, C. Marquet, M. McPartlon, T. Geffner, E. Rossi, G. Corso, H. Stärk, Z. Carpenter, E. Kucukbenli, M. Bronstein, and L. Naef. PINDER: The protein interaction dataset and evaluation resource. *bioRxiv preprint*, 2024. doi: 10.1101/2024.07.17.603980.
- K. Lee, D. Ippolito, A. Nystrom, C. Zhang, D. Eck, C. Callison-Burch, and N. Carlini. Deduplicating training data makes language models better. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8424–8445, 2022. doi: 10.18653/v1/2022.acl-long.577.
- W. Li and A. Godzik. Cd-hit: a fast program for clustering and comparing large sets of protein or nucleotide sequences. *Bioinformatics*, 22(13):1658–1659, 2006. doi: 10.1093/bioinformatics/btl158.
- I. Magar and R. Schwartz. Data contamination: From memorization to exploitation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 157–165, 2022. doi: 10.18653/v1/2022.acl-short.18.
- F. I. Marin, F. Teufel, M. Horlacher, D. Madsen, D. Pultz, O. Winther, and W. Boomsma. BEND: Benchmarking DNA language models on biologically meaningful tasks. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024. doi: 10.48550/arXiv.2311.12570.
- E. Nguyen, M. Poli, M. Faizi, A. Thomas, C. Birch-Sykes, M. Wornow, A. Patel, C. Rabideau, S. Massaroli, Y. Bengio, S. Ermon, S. A. Baccus, and C. Ré. HyenaDNA: Long-range genomic sequence modeling at single nucleotide resolution. *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. doi: 10.48550/arXiv.2306.15794.
- C. G. Northcutt, A. Athalye, and J. Mueller. Pervasive label errors in test sets destabilize machine learning benchmarks. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2021. doi: 10.48550/arXiv.2103.14749.
- A. Patel, A. Singhal, A. Wang, A. Pampari, M. Kasowski, and A. Kundaje. DART-Eval: A comprehensive DNA language model evaluation benchmark on regulatory DNA. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, pages 6204–62061, 2024. doi: 10.52202/079017-1981.
- S. Petti and S. R. Eddy. Constructing benchmark test sets for biological sequence analysis using independent set algorithms. *PLOS Computational Biology*, 18(3):e1009492, 2022. doi: 10.1371/journal.pcbi.1009492.
- A. M. Rafi, B. Kiyota, N. Yachie, and C. de Boer. Characterizing homology-induced data leakage and memorization in genome-trained sequence models. *bioRxiv preprint*, 2026. doi: 10.1101/2025.01.22.634321. URL <https://www.biorxiv.org/content/10.1101/2025.01.22.634321v2>. version 2, posted 25 May 2026.
- O. Sainz, J. Campos, I. García-Ferrero, J. Etxaniz, O. Lopez de Lacalle, and E. Agirre. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10776–10787, 2023. doi: 10.18653/v1/2023.findings-emnlp.722.
- J. Schreiber, R. Singh, J. Bilmes, and W. S. Noble. A pitfall for machine learning methods aiming to predict across cell types. *Genome Biology*, 21(1):282, 2020. doi: 10.1186/s13059-020-02177-y.

- M. Steinegger and J. Söding. MMseqs2 enables sensitive protein sequence searching for the analysis of massive data sets. *Nature Biotechnology*, 35(11):1026–1028, 2017. doi: 10.1038/nbt.3988.
- F. Teufel, M. H. Gíslason, J. J. Almagro Armenteros, A. R. Johansen, O. Winther, and H. Nielsen. GraphPart: homology partitioning for biological sequence analysis. *NAR Genomics and Bioinformatics*, 5(4):lqad088, 2023. doi: 10.1093/nargab/lqad088.
- S. Toneyan, Z. Tang, and P. K. Koo. Evaluating deep learning for predicting epigenomic profiles. *Nature Machine Intelligence*, 4(12):1088–1100, 2022. doi: 10.1038/s42256-022-00570-9.
- I. Walsh, G. Pollastri, and S. C. E. Tosatto. Correct machine learning on protein sequences: a peer-reviewing perspective. *Briefings in Bioinformatics*, 17(5):831–840, 2016. doi: 10.1093/bib/bbv082.
- S. Whalen, J. Schreiber, W. S. Noble, and K. S. Pollard. Navigating the pitfalls of applying machine learning in genomics. *Nature Reviews Genetics*, 23(3):169–181, 2022. doi: 10.1038/s41576-021-00434-9.
- D. R. Zerbino, S. P. Wilder, N. Johnson, T. Juettemann, and P. R. Flicek. The Ensembl Regulatory Build. *Genome Biology*, 16:56, 2015. doi: 10.1186/s13059-015-0621-5.
- J. Zhou and O. G. Troyanskaya. Predicting effects of noncoding variants with deep learning-based sequence model. *Nature Methods*, 12(10):931–934, 2015. doi: 10.1038/nmeth.3547.
- K. Zhou, Y. Zhu, Z. Chen, W. Chen, W. X. Zhao, X. Chen, Y. Lin, J.-R. Wen, and J. Han. Don't make your LLM an evaluation benchmark cheater. *arXiv preprint*, 2023a. doi: 10.48550/arXiv.2311.01964.
- Z. Zhou, Y. Ji, W. Li, P. Dutta, R. Davuluri, and H. Liu. DNABERT-2: Efficient foundation model and benchmark for multi-species genome. *arXiv preprint*, 2023b. doi: 10.48550/arXiv.2306.15006.